

# Booting and System Management Daemons

## Unix – Chapter 2

Bjarte Wang-Kileng

HVL

January 15, 2025



**Western Norway  
University of  
Applied Sciences**

# Boot process and configuration

## Select boot device:

- ▶ Enable disk as boot device.
- ▶ Choose boot device.
- ▶ Specify boot device priorities.


## Determine which kernel to boot:

- ▶ Select what kernel to load.
- ▶ Specify default kernel.


## Execute startup scripts:

- ▶ Specify what init scripts to run.
- ▶ Run order or dependencies between init scripts.

Linux & UNIX boot process



- ▶ The acronyms:
  - BIOS = *Basic Input/Output System*
  - UEFI = *Unified Extensible Firmware Interface*
- ▶ BIOS and UEFI are firmware that is run when the computer starts.
  - CPU is hardwired to run the firmware.
- ▶ Usually stored in **EEPROM** on the motherboard.
- ▶ Typically includes a GUI to e.g. select the boot disk, or specify a disk priority for boot.
- ▶ Initializes the computer and runs the next stage of the bootstrapping code (e.g. GRUB).
- ▶ UEFI has replaced BIOS on modern PCs.
- ▶ UEFI firmwares usually implement some kind of BIOS compatibility.
  - UEFI can boot from a MBR (MBR on next slide).



- ▶ Boot disk must start with MBR (*Master Boot Record*) that contains:
  - First-stage boot loader, *boot block*.
  - Disk partition table.
- ▶ The second-stage boot loader, e.g. GRUB:
  - Includes the necessary file system drivers to load the OS.
  - Loads the OS.
- ▶ The second-stage boot loader is read either from:
  - The active partition, or
  - the dead zone before the first partition (64 disk blocks),
- ▶ All BIOS configurations are done inside the firmware.
  - BIOS has no concept of OS or file systems.

- ▶ UEFI = Unified Extensible Firmware Interface
- ▶ GUID = Globally Unique Identifier
- ▶ GPT = GUID Partition Table
- ▶ ESP = EFI System Partition
- ▶ FAT = File Allocation Table (MS\_DOS file system)
- ▶ NVRAM = Non-Volatile RAM

- ▶ UEFI firmware can handle FAT partitions.
- ▶ ESP is a FAT partition that stores the target application, e.g. GRUB.
  - Usually mounted as “/boot/efi”.
- ▶ UEFI parameters are stored in NVRAM, accessible by UEFI and OS.
  - Accessible as “/sys/firmware/efi/efivars”, or “/sys/firmware/efi/vars”.
- ▶ GPT identifies the ESP and the target application.
  - Stored as UEFI parameters.
  - Default target application is “/efi/boot/bootx64.efi”.
- ▶ UEFI configurations can be modified in user space using e.g. the command **efibootmgr** or through the “/sys/firmware/” file system.

- ▶ Default boot loader for Linux.
- ▶ Can boot multiple operating systems.
- ▶ GRUB code location:
  - UEFI – ESP
    - E.g. “EFI/almalinux/grubx64.efi”.
  - BIOS – Usually in the dead zone before the first partition.
- ▶ GRUB can read its configurations from a file.
  - Usually found as **grub.cfg** in the ESP or “/boot/grub2” partition.

# Modify GRUB

- ▶ Modifications of **grub.cfg** do not survive updates.
  - Configuration file “/etc/default/grub”.
  - Command `grub2-editenv` – modify GRUB parameters
- ▶ Transfer the changes to the **grub.cfg** file with the command:

```
grub2-mkconfig -o <path-to>/grub.cfg
```

- Before using *grub2-mkconfig*, package “os-prober” must be installed.
- ▶ To specify a default kernel, see e.g. [GRUB 2 - Fedora Project Wiki](#).



- ▶ Lets us modify the boot entries at boot time.
  - Changes are not saved.
- ▶ Modify a GRUB menu entry:
  - Locate the GRUB menu entry.
  - Press the “e” key.
  - Modify the kernel line.
  - Use “F-10” to boot from the modified GRUB stanza.
- ▶ See the book or Internet for details.

- ▶ Started by the kernel as process with pid equal to 1.
- ▶ Starts and stops system services and daemons.
- ▶ The systemd process is the top process in the process hierarchy.

- ▶ Most major distros uses systemd to initialize the system.
  - RedHat did use the init system of System V before RedHat 7.
- ▶ Startup sequence:
  - System V init – the scripts are run in a set sequence.
  - systemd – dependencies determine the startup sequence.
- ▶ Parallelism:
  - System V init – the scripts are run in a strict sequence, one at a time.
  - systemd – can run the startup scripts in parallel.
- ▶ systemd is a more complicated system.

# systemd units and unit files

- ▶ A unit is an entity that is managed by systemd:
  - Can be a service started by a script.
  - Can be a target (more later).
  - And much more.
- ▶ A systemd unit is described by a unit file:
  - Path of executable.
  - Specify how to start and stop.
  - Specify dependencies.
- ▶ Unit files are read from:
  - “/etc/systemd/system” – Highest priority.
  - “/usr/local/lib/systemd/system”
  - “/usr/lib/systemd/system”
  - “/lib/systemd/system”
- ▶ Unit-file suffix specify unit type:
  - E.g. “.service”, “.target”, “.timer”, “.socket”, “.mount”, . . . .

# Find systemd units

- ▶ List all loaded units:

```
systemctl list-units
```

- ▶ List loaded services only:

```
systemctl list-units --type=service
```

- ▶ List all service unit files:

```
systemctl list-unit-files --type=service
```

# Enable/disable/mask

- ▶ An enabled service will be started at “boot” by systemd.

```
systemctl enable mariadb.service
```

- ▶ A disabled service will not be started at boot by systemd.
  - Can still be started manually.
  - Can also be started due to a depending enabled service.

```
systemctl disable mariadb.service
```

- ▶ A masked service can not be started.
  - Unit file is linked to “/dev/null”.

```
systemctl mask mariadb.service
```

## Enabled unit and boot

An enabled unit is started if any of the *WantedBy* units are started. Usually, this means that the unit is started at boot.

# Start/stop/status

- ▶ A service can be started.

```
systemctl start mariadb.service
```

- ▶ A service can be stopped.

```
systemctl stop mariadb.service
```

- ▶ We can check the status of a service.

```
systemctl status -l mariadb.service
```

- ▶ States are e.g. *enabled*, *disabled* or *static*.
- ▶ Only unit files with an `Install` section can be *enabled* or *disabled*.
- ▶ The **list-unit-files** sub command list service and enablement state.
  - See the man page for all the enablement states.
- ▶ Services with state *static* have no `Install` section.
  - Can only be started by hand, or if named as a dependency by another service.



# Targets

- ▶ Targets are used to gather units.
  - Action on target will act on all units of the target.

- ▶ List loaded targets:

```
systemctl list-units --type=target
```

- ▶ List all targets:

```
systemctl list-unit-files --type=target
```

- ▶ Some targets have special meaning.
  - Run automatically by systemd at specific events.

# Some special targets

`ctrl-alt-del.target`: Run when Control+Alt+Del is pressed.

`default.target`: Default unit to start at boot.

`emergency.target`: Run if boot fails, e.g. due to a failing local disk mount.

▶ Play with this mode **before** you need it.

`graphical.target`: Sets up the graphical login screen.

`multi-user.target`: Sets up a multi-user system.

`reboot.target`: Shutdown and reboot the system.

`shutdown.target`: Shutdown the system.

`rescue.target`: Pulls in the base system for administrative purposes.

# Using targets

- ▶ The sub command **isolate** changes the current mode:

```
systemctl isolate multi-user.target
```

- Activate the target with its dependencies, and deactivate all others.

- ▶ Show the default target:

```
systemctl get-default
```

- ▶ Set the default target:

```
systemctl set-default multi-user.target
```

- ▶ Can also specify a target on the kernel boot line
  - Emergency mode: `systemd.unit=emergency.target`
    - Short form: `emergency`
  - Rescue mode: `systemd.unit=rescue.target`.
    - Short form: `rescue`

# Unit dependencies

- ▶ Dependencies are specified in the unit files.
- ▶ Keywords are e.g. *Wants*, *Requires*, *Requisite*, *Bindsto*, *PartOf*, *Conflicts*
  - Does not imply any sequence for processing.
  - See the book page 51 and manual pages for their meaning.
- ▶ *Wants* are a weaker form of *Requires*.
  - Success of unit does not depend on success of *Wants*.
  - Success of unit does depend on success of *Requires*.
- ▶ Keywords to serialize processing are e.g. *After*, *Before*.
  - If no serializing keyword, systemd will try to run processes in parallel.
- ▶ Systemd assumes a default set of dependencies for units.
  - Turn off assumption with `DefaultDependencies=false`

- ▶ New unit files can be created in an text editor.
- ▶ The systemctl sub command **edit** lets us modify an existing unit-file.

```
systemctl edit mariadb.service
```

- Creates an override in “/etc/systemd/system/mariadb.service.d/”.
- ▶ Do not modify existing unit-files.
  - Use an override, e.g.  
“/etc/systemd/system/mariadb.service.d/override.conf”.
- ▶ Using unit-file overrides:
  - The original and the override will be merged by systemd at use.
  - If collisions, the override has higher priority.
- ▶ Install section can not be modified by an override.

# Show unit properties

- ▶ List unit, with overrides:

```
systemctl cat mariadb.service
```

- ▶ List specific property of unit:

```
systemctl show -p After mariadb.service
```

# Some comments on logging

- ▶ Systemd has its own logging system, managed by the *journald* daemon.
- ▶ Systemd log messages are stored in the “/run” directory.
  - Typically, the rsyslog daemon will process also the systemd logs.
- ▶ Can be accessed with the command **journalctl**.
  - Show log of the Bluetooth unit:

```
journalctl -u bluetooth.service
```

- ▶ System logging will be covered later.