

# Unix – Chapters 3, 4, 8

Bjarte Wang-Kileng

HVL

January 16, 2025



**Western Norway  
University of  
Applied Sciences**

# Outline

- 1 Chapter 3 – Access Control and Rootly Powers
- 2 Chapter 4 – Process Control
- 3 Chapter 8 – User Management

# Outline

- 1 Chapter 3 – Access Control and Rootly Powers
- 2 Chapter 4 – Process Control
- 3 Chapter 8 – User Management

# Standard UNIX Access control

- ▶ Login as a user.
- ▶ Users are member in one or several groups of users.
- ▶ Processes and files have owners.
- ▶ User *root* has special privileges.
- ▶ Linux kernel allows for user namespaces.
  - Not all of this chapter apply if using user namespaces.

# UID and GID

- ▶ System track users and groups by a numeric id.
- ▶ User names are mapped to UIDs through the file “/etc/passwd”.
- ▶ Group names are mapped to GIDs through the file “/etc/group”.

# File system access control

- ▶ A file belongs to one owner and one group.
- ▶ File owner can set permissions on file.
- ▶ File owner can specify group of file.
  - File owner must be member of this group.
  - User *root* can assign any group to a file.
- ▶ Access to file can be specified for owner, group and others.

# Process ownership

**RUID and RGID:** Real UID and Real GID

The real owner. Can send signals to the process and change scheduling policy.

**EUID and EGID:** Effective UID and Effective GID.

Used for access checks.

Needed since processes can be run with privileges different from that of the owner.

**SUID and SGID:** Saved UID and Saved GID.

Needed when a process with elevated EUID temporarily change its EUID to RUID.

**FUID and FGID:** File system UID and File system GID.

Used for access control to the file system. Usually equal to EUID.

# The root account

- ▶ Has UID and GID equal to 0.
- ▶ Can perform valid operations on any file or process.
  - Not necessarily true on a network filesystem
- ▶ Processes run as root can change its UID and GID.



# setuid and setgid

- ▶ Programs with setuid or setgid permissions set will run with EUID or EGID equal to that of the program file.
- ▶ Used by e.g. the **passwd** program.
- ▶ Using capabilities (later), subset of root permissions can be given to programs.

# Issue commands as root

- ▶ Login as root.
- ▶ Using command **su**.
- ▶ Using command **sudo**.

# Login as root

- ▶ No log or record of operations.
- ▶ Not recommended for a production system.
- ▶ Can disable the root account:

```
passwd -l root
```

- Observe that the emergency and rescue boot modes require root login.

```
You are in rescue mode. After logging in, type "journalctl -xb" to view
system logs, "systemctl reboot" to reboot, "systemctl default" or "exit"
to boot into default mode.
```

```
Cannot open access to console, the root account is locked.
See sulogin(8) man page for more details.
```

```
Press Enter to continue.
```

- ▶ Enable root login:

```
sudo passwd -u root
```

# Using command **su**

- ▶ Log entry of who run the **su** command, but not of the operations.
- ▶ Switch **--login** gives an environment similar to a real login.

```
su --login # or su -
```

- ▶ Root can **su** to any user without a password.

```
su --login username
```

# Using command **sudo**

- ▶ Run one command as root (or another user).
- ▶ Creates a log entry of command and user.
- ▶ Configurable through the file “/etc/sudoers” or directory “/etc/sudoers.d”.
  - Use the command **visudo** to edit “/etc/sudoers”.
  - If several lines apply, the last matching line applies.
- ▶ Can give a user root access to perform specific tasks.
- ▶ Switch *--login* or *-i* gives an environment similar to a real login.
  - Similar to **su** with switch *--login*.
- ▶ Switch *--list* or *-l* to list sudo rights of invoking user.

# Unattended use of sudo

- ▶ E.g. cron jobs.
- ▶ Must use the *NOPASSWORD:* parameter.

```
%MYSQL_ADMINS ALL = (mysql) NOPASSWORD: /usr/local/bin/mysqlbackup
```

- Only use *NOPASSWORD* for specific commands.

- ▶ Must allow in sudoers file for sudo without a terminal.

```
Defaults !requiretty
```

Should be OK by default configuration.

# System accounts

- ▶ Systems typically include many system users and groups.
- ▶ Uses low values for UID and GID.
  - See the file “/etc/login.defs”.
- ▶ Usually no login.
  - Shell set to “/bin/nologin”
- ▶ Usually have no elongated access rights, but own files and processes.
  - MariaDB runs as user and group *mysql*.
  - Directory structure “/var/lib/mysql” belongs to a user and group *mysql*.

# PAM (Pluggable Authentication Module)

- ▶ Authentication framework used on modern Linux systems.
- ▶ Pluggable system for user authentication.
- ▶ See e.g. the manual page **pam.conf(5)**.
- ▶ More later.



# File system access lists (ACL)

- ▶ Generalization of the user/group/other permissions.
- ▶ Can set file system permissions for specific users and groups.
- ▶ Part of the file system – Supported by all major UNIX and Linux file systems.

# Linux Capabilities

- ▶ Traditionally, Linux authorization uses two levels only:
  - Full root access.
  - Normal user access.
- ▶ Linux Capabilities divide root powers into approximately 30 separate permissions.
- ▶ For a overview of all capabilities, see e.g. the manual **capabilities(7)**.
- ▶ Capabilities can be given to program files.
  - Similar to **chmod u+s**, but with less privileges.
  - Processes started from the program file gets the capabilities.
- ▶ Capabilities can be given to processes.
  - Process can start child processes and give of its capabilities.

# Linux namespaces

- ▶ System that wrap global resources into sandboxed environments.
- ▶ Isolated containers for resources.
- ▶ Users and processes only see the sandboxed environment.
- ▶ Mount namespace is an exception:
  - Namespace see the mounts of its parent, but
  - the parent does not see the mounts of child mount namespaces.
  - Can together with **chroot** create isolated islands of file systems.

# The Linux namespaces

- ▶ PID (processes),
- ▶ network (network interfaces, routing),
- ▶ UTS (hostname),
- ▶ user (UIDs),
- ▶ mount (mount points, file systems),
- ▶ IPC (System V IPC).

# LSM – Linux Security Modules API

- ▶ API that allow security modules as loadable kernel modules.
- ▶ E.g. SELinux, AppArmor, Smack, TOMOYO, Yama.
- ▶ Current LSMs do not cooperate – Use only one.

# MAC – Mandatory Access Control

- ▶ Access control policies supplement or override traditional model, e.g.:
  - Web documents must reside in “/var/www/html”
  - Only the user can access his home directory.
  
- ▶ Both SELinux and AppArmor are MAC systems.

# SELinux

- ▶ Created by NSA – Open source.
- ▶ Used by Fedora and Redhat with clones, but available also for other distributions.
- ▶ Both MAC and role based access control.
- ▶ Level of control can be set in file **/etc/selinux/config**.
- ▶ Processes are monitored, and only actions consistent with the policies are allowed.
- ▶ Security context of files are stored as extended attributes.
- ▶ Security context is compared with policies encoded in policy files.
- ▶ Details in lab assignment.

# AppArmor

- ▶ By Canonical (Ubuntu).
- ▶ MAC system to supplement the traditional access control system.
- ▶ Goal is to limit damage from services.
- ▶ Action must be allowed by both the traditional model and AppArmor.



# Outline

- 1 Chapter 3 – Access Control and Rootly Powers
- 2 Chapter 4 – Process Control
- 3 Chapter 8 – User Management

# Process parameters

- ▶ PID: Unique number that identifies the process.
  - Uniqueness only within the same PID namespace.
- ▶ PPID: The parent process.
  - Processes are created in a tree structure.
- ▶ RUID, EUID, RGID, EGID: See chapter 3.
- ▶ Niceness: Scheduling parameter.
  - Low nice values give more CPU time.
- ▶ Control terminal: Terminal of STDIN, STDOUT and STDERR.

# Signals to processes

- ▶ The command **kill** can send signal to a process.

```
kill [-s signal] pid
```

- ▶ See book and man page **signal(7)** for more signals:
  - SIGTERM (15): Default signal sent by kill. Terminates the process.
  - SIGKILL (9): Signal can not be caught. Process is killed.
  - SIGINT (2): Signal sent by Ctrl-C.
  - SIGSTOP (19): Freeze process.
  - SIGCONT (18): Continue process stopped with SIGSTOP.
- ▶ Always try SIGTERM (default) before SIGKILL.
- ▶ Sometimes even SIGKILL will fail (e.g. disk problem).
  - Only reboot will remove process.
- ▶ Also **pkill** and **killall**:

```
killall httpd # All httpd processes  
pkill -u bki,jon emacs # All emacs processes of bki and jon
```

# Processes commands

► Monitoring – **ps**, **pidof**, **pgrep**.

```
ps -C emacs  
pidof /usr/bin/emacs  
pgrep -u bki emacs
```

► Interactive monitoring – **top**.

- “q” to quit, “?” for help.

► Change the nice value – **nice**, **renice**.

- Non-root user can only increase the value.
  - Non-root user can not reset the value.
- Root can set the value arbitrarily.

► Trace system calls and signals – **strace**.

# Process scheduling in Linux

- ▶ Uses multi-level queue scheduling with preemption – 101 queues.
- ▶ Kernel uses two different scheduling systems:
  - 100 real time queues – Multi-level queue scheduling.
  - One queue for user- and interactive processes – Organized as a self balancing binary tree (a *red-black tree*).
- ▶ The 101 queues have priorities from 0 (lowest) to 100 (highest).
  - A process belonging to a higher priority queue will preempt a running lower priority process.

# The real time queues

- ▶ Uses 100 real time queues with priorities from 1 to 100.
- ▶ Within each queue a process can belong to scheduling class **SCHED\_FIFO** or **SCHED\_RR**.
- ▶ A **SCHED\_FIFO** process is not preempted unless a higher priority process needs the CPU.
- ▶ **SCHED\_RR** uses a time quantum of 100ms.

# User- and interactive processes

- ▶ The queue of user- and interactive processes has priority 0 (lowest).
- ▶ Processes usually belong to scheduling class **SCHED\_OTHER**, but can also be **SCHED\_IDLE** or **SCHED\_BATCH**.
- ▶ The importance of a process is determined by its scheduling class and a value called the *nice* value.
  - *nice* values go from -20 (most favorable) to +19 (least favorable).
- ▶ The scheduling algorithm is named **Earliest Eligible Virtual Deadline first** (EEVDF).
- ▶ EEVDF improves on the older *Completely Fair Scheduler* (CFS):  
*CFS always tries to split up CPU time between runnable tasks as close to "ideal multitasking hardware" as possible.*

# CFS and the *nice* value

- ▶ Each process is assigned a “virtual” run time.
- ▶ The “virtual” run time of a process increases monotonic with its CPU time, normalized to the minimum run time of the queue.
  - [How does the Completely Fair Scheduler prevent starvation . . .](#)
- ▶ The “virtual” run time clock ticks slower with lower *nice* values.
- ▶ The “virtual” run time clock ticks slower for processes of **SCHED\_OTHER** compared to processes of **SCHED\_IDLE** and **SCHED\_BATCH**.
- ▶ CFS will run the process with the shortest “virtual” run time.
  - Low *nice* values give more CPU time.



# The “/proc” file system

- ▶ Linux uses several pseudo file systems with “files” that are dynamically created by the kernel.
  - Not real files, but
  - entry points to kernel data.
- ▶ All process information is stored in “/proc”, e.g. info on *systemd* below “/proc/1/”.
- ▶ Commands like **ps** and **top** get their information from “/proc”.

# Periodic processes

- ▶ Cron
- ▶ Anacron
- ▶ Systemd timers

# Cron

- ▶ Crontab files configure tasks to be run at certain times.
- ▶ Main crontab file – “/etc/crontab”
- ▶ Scripts in directories “/etc/cron.d”, “/etc/cron.hourly”, “/etc/cron.daily”, “/etc/cron.weekly”, “/etc/cron.monthly”.
- ▶ Users must use command **crontab** to create crontab file.
  - Users crontab files in “/var/spool/cron”.
- ▶ For format of crontab file:

```
man 5 crontab
```

- ▶ Logging:
  - If syslog, in file “/var/log/cron”.
  - Systemd journal, e.g.:

```
journalctl --unit=crond.service
```

# Anacron

- ▶ Similar to Cron, but does not assume machine to be on at all times.
  - Will run missed job when computer is turned on.
- ▶ Main configuration file – “/etc/anacrontab”
- ▶ Cron and Anacron are working together on RedHat 9 and clones.
  - The *crond.service* is running *anacron* one minute past every hour, then
  - Anacron is running the daily, weekly and monthly Cron jobs.

# systemd timers

- ▶ Absolute and relative times, e.g. 30 seconds after boot:

```
OnBootSec=30
```

- See book for more parameters.
- ▶ Can be started and enabled as other systemd units.

# Outline

- 1 Chapter 3 – Access Control and Rootly Powers
- 2 Chapter 4 – Process Control
- 3 Chapter 8 – User Management

# Account mechanism

- ▶ Login information can have many sources, see the file “/etc/nsswitch.conf”.
  - E.g. local files, SSSD, NIS.
  - The SSSD daemon can use eg. LDAP, Kerberos and Active Directory for authentication.
    - More details later and in lab assignment.
  
- ▶ Traditionally, account information is found in the file “/etc/passwd”.

# Files

“/etc/passwd”: User name, password, uid, gid, gecoss, home, shell.

- ▶ GECOS – Readable personal info.
- ▶ Password usually in shadow file.

“/etc/shadow”: User name, encrypted password, password lifetime fields.

“/etc/group”: Group name, password, gid, list of member uids.

- ▶ Group password is rarely used.
- ▶ Group password can allow non members to enter group.

“/etc/gshadow”: As “/etc/shadow”, but for groups.

- ▶ Rarely used.



# Passwords

► Stored encrypted:

```
$prefix$options$salt$password-hash
```

- The prefix field specifies the hash method.
  - Value **6** specifies SHA-512, and was default on e.g. RHEL 7 and 8.
    - No *options* field.
  - Value **y** specifies [yescrypt](#), and is default on RHEL 9.
- Configuration of password strength and number of login attempts through PAM modules.
- E.g. **pam\_pwquality**, **pam\_faillock**.

# Working with users and groups

- ▶ GUI programs.
  - Not suitable for bulk processing and scripting.
- ▶ Using commands.
- ▶ Working with the account system files.
  - The commands **vipw** and **vigr** will lock files before opening an editor.

# Commands

User commands: **passwd**, **useradd**, **usermod**, **userdel**, **chfn**, **chage**, **chsh**, **newusers**.

▶ **passwd**, **chfn** and **chsh** also by normal users.

Group commands: **gpaswd**, **groupadd**, **groupmod**, **groupdel**.

▶ Command **gpaswd** also by group administrator.

Information: **getent**, **id**, **whoami**, **finger**, **pinky**.

▶ Also normal users.

Work as another: **su**, **sg**, **runuser**, **newgrp**, **sudo**.

Access: **chown**, **chmod**, **setfacl**, **chgrp**.

# Configuration files

“/etc/default/useradd”: Defaults for **useradd**.

“/etc/skel”: Content to fill in HOME with **useradd**.

“/etc/login.defs”: Shadow password suite configuration.

“/etc/security”: Configuration files for login.

# Remove users

- ▶ Accounts can be locked and unlocked.
  - `usermod -L` and `usermod -U`.
- ▶ After removing user, e.g. using `userdel`, check for remnant files.

```
find filesystem -xdev -nouser
```

- ▶ Remember also databases, phone lists, crontab, pending jobs, processes, mail spool etc.