

# The Linux Filesystem

Chapter 5 in Unix and Linux System Administration Handbook

Haakon André Reme-Ness


HVL

`harn@hvl.no`

January 29, 2025

# The Windows filesystem


Before jumping to the Linux filesystem, let's first have a look at one that “everyone” is familiar with: the **Windows** filesystem.



# The Windows filesystem


- ▶ Divided into drives
  - **A & B** reserved for floppy drives
  - **C** as the system boot drive
  - Additional drives labelled **D and beyond**
    - DVD drives
    - External HDDs
    - Flash drives

# The Windows filesystem



# The Linux filesystem

**Linux** (and other Unix-like systems) organise things differently.



# The Linux filesystem



# The Linux filesystem

- ▶ The filesystem is presented as a single unified hierarchy
  - Starts at /, the **root directory**
- ▶ Continues downwards through an arbitrary number of subdirectories
  - Directory = “folder”
- ▶ Can be arbitrarily deep
  - Certain limitations apply

# The Linux filesystem


- ▶ Contains a *representation* of files (as expected) ... but also **a lot more!**
  - Processes
  - Audio devices
  - Kernel data structures and parameters
  - Interprocess communications channels
- ▶ In Linux (and Unix generally), **everything is a file**
- ▶ Why?
  - It's convenient!
  - Consistent APIs and easy access from shell
  - Comes with the disadvantage of having a patchwork of multiple different filesystem implementations




# Filesystem mounting

- ▶ The filesystem is composed of smaller chunks containing a **directory** and its **subdirectories**
  - Each directory is a **branch** in the **file tree** starting at the root `/`
  - These directories are also known as *filesystems*
- ▶ With there being only one “drive” in Unix (the root directory `/`) how are additional drives and partitions represented?
  - Filesystems also live on disk partitions and other physical volumes
  - And these can be attached to the file tree using the **mount** command

# Filesystem mounting



# Filesystem mounting




# Filesystem mounting

- ▶ **mount** maps a directory within the file tree, called the **mount point**, to the root of another filesystem
- ▶ The mount point is often an empty directory
  - e.g., /dev/sda in the previous example
  - Contains the drive filesystem after mounting
- ▶ But the mount point can also already contain files
  - But previous files in that directory will become unavailable while the new filesystem is mounted

# Examples using mount

- ▶ Mount “/dev/sda” to “/users”  
**sudo mount /dev/sda /users**
- ▶ Unmount the previous directory  
**sudo umount /users** (-f and -l flags available)
- ▶ List all mounted filesystems  
**mount**

# File tree organisation




- ▶ UNIX systems have never been well organised
  - Multiple incompatible naming conventions used simultaneously
  - Files scattered randomly around the namespace
  - Files can be grouped by function, and not by how often they are likely to change
    - Makes upgrades harder
  - No unified folder for application files
    - e.g., such as “Program Files” on Window

- ▶ So, when installing new software, trust the installer!
- ▶ Do not change default location, unless you have a very compelling reason to do so
- ▶ While it might be tempting to reorganise certain files, this will likely only create problems
  - Hidden dependencies



- ▶ Despite the apparent chaos, some directories are worth mentioning
  - /home – User directories
  - /etc – Critical system and configuration files
  - /tmp – Temporary files
  - /usr – Standard non-critical programs
  - /var – Spool directories, accounting information and log files
  - /dev – Devices (disks, printers & the like)



- ▶ Most implementations of UNIX filesystems define seven types of files:
  - Regular files
  - Directories
  - Character device files
  - Block device files
  - Local domain sockets
  - Named pipes
  - Symbolic links
- ▶ Everything listed in the file system must be constrained to one of these file types
- ▶ This also applies for processes, audio devices, etc

## ► Regular files

- Consists of a series of bytes of any structure
- Can be anything from text files, data files, executable files to library files

## ► Directories

- Contains named references to other files
  - Can be created with `mkdir` and removed with `rmdir` (if empty)
    - Non empty directories can only be deleted with `rm -r`
  - All directories also contain the entries `.` and `..`
    - `.` refers to the directory itself
    - `..` refers to the parent directory
- Note:** `..` points to itself in the root directory `/`

## ► Hard links

- The name of a file is stored within its parent directory, and not within the file itself
- This allows more than one directory that points to the same file **inode**
  - Each **file** is associated to an integer index which is the inode numbers of the files
  - Every UNIX file has a **unique inode**
- This creates the illusion that the file exists multiple places
- The file will not be deleted until all these pointers (links) are deleted

## ► Symbolic links

- Also known as a “soft” link
- A symbolic link is a file that having its own inode number
- Unlike a hard link, which is a direct reference to the file, a symbolic links is a reference to a filename

# The Linux filesystem




Figure: Visualised difference between hard links and symbolic links



## ► Character and block device files

- Only a slight distinction between these two types
- Both used for communications with the system's hardware and peripherals
- Look like regular files, but passes requests onwards to the device driver when called.
- Allows the kernel to be relatively abstract and hardware independent
- Characterised by two numbers
  - Major device number: tells the kernel what driver the file refers to
  - Minor device number: tells the driver which physical unit to address

## ► Local domain sockets

- Connections between processes that allows for communications
- Local domain sockets are only accessible from the local host, and are referred to through a filesystem object
  - Unlike network sockets, that communicate through network ports

## ► Named pipes

- Similar to local domain sockets, allowing processes on the same host to communicate
- Also known as FIFO pipes (i.e., first-in, first-out)
- The existence of both local domain sockets and named pipes, despite having similar purposes, is due to historical reasons

# File attributes

Bit:	12	11	10	9	8	7	6	5	4	3	2	1
	SUID	SGID	Sticky	r	w	x	r	w	x	r	w	x
	Bits for special access privileges			Owner's access privileges			Group's access privileges			Others' access privileges		

- ▶ Under the traditional UNIX/Linux filesystem model, every file has
  - Nine permission bits
  - Three other bits that constitute a file's "mode"
    - Affects the operation of executable files
  - Four bits of file type information
    - Set during file creation, and can not be changed later
    - (not shown in above figure)

## ► The permission bits

- Each file has **nine**
- Determine what operations can be performed on a file, and by whom
- Divided into three sets, to define access for
  - The file owner
  - The file owner group
  - Everyone else
- Each of these sets has three bits
  - A read bit (**r**)
  - A write bit (**w**)
  - An execute bit (**x**)

Example, a file with the following nine permissions bits:

**111 110 100**

**111**: the owner has read, write and execute permissions

**110**: the group has read and write permissions

**100**: Everyone else have read permissions only

► Octal numbers

- The three bits of each set (read, write, execute) are often written in terms of octal numbers (0 – 7)

# File attributes

Octal	Binary	Perms	Octal	Binary	Perms
0	000	---	4	100	r--
1	001	--x	5	101	r-x
2	010	-w-	6	110	rw-
3	011	-wx	7	111	rwX

Figure: Conversion table between octal and binary numbers

For example, the following permission bits for a file:

**111 110 100**

Can be written as in octal numbers:

**7 6 4**



- ▶ The permission bits of a file can be changed using the **chmod** command
  - These can be specified using octal numbers. E.g.,

## **chmod 764**

the permissions of a file to read, write and execute for the owner, read write for the group, and read only for all others:

- Alternatively, a mnemonic syntax is also accepted. It combines a set of targets (**u**, **g**, **o** for **u**ser, **g**roup and **o**thers, or **a** for all three) combined with an operator (+, -, = for add, remove or set) before providing the wanted permissions. E.g.,

## **chmod go+rw**

To add **r**ead and **w**rite permissions to the **g**roup and all **o**thers

- ▶ With the mnemonic syntax being available, why then bother with octals?
  - Preferred by many for being shorter and easier to type.
  - Many tools only accept octal numbers, and may respond with them when queried for information.

In addition to the nine permission bits, there are also **three other bits**

- ▶ **The setuid and setgid bits** (octal value: 4000 & 2000)
  - When on **executable files**:  
allows programs to access files and processes that would otherwise be off-limits to the current user
  - When set on a **directory**:  
setgid allows newly created files here to have the same group ownership as the directory, instead of the user who created the file
    - Simplifies sharing

- ▶ **The sticky bit** (octal value: 1000)
  - When set on a **directory**: prevents a file/directory from being deleted or modified when set, unless you're the owner or the superuser.
  - Mostly obsolete, and ignored by most Linux distributions for regular files. Still occasionally used on folders.

# Access Control Lists (ACLs)

- ▶ A more powerful (**but also more complicated**) way to handle file permissions
- ▶ Each file or directory can have an associated ACL that lists the permission rules to be applied to it
- ▶ No set length, and can contain permission specifications for multiple users or groups
- ▶ Allows specifying partial permissions, negative permissions and inheritance features
  - Allows access specifications to be propagated to newly created filesystem entities

# Access Control Lists (ACLs)

- ▶ However
  - Exists mainly to serve a certain niche that lies outside the mainstream of UNIX and Linux administration
  - Primarily to facilitate Windows compatibility
  - Though some enterprises may also require the added flexibility
- ▶ May be more trouble than it is worth
  - Tedious to use
  - Can cause problems when communicating with systems not using ACL
  - Tends to become increasingly unmaintainable